
‘Spark project’

Release 0.0.1

Bram Vonk

Dec 02, 2021

CONTENTS:

1 General	3
1.1 About JADS	3
1.2 About Enexis	3
1.3 About the current situation	4
1.4 About this project's aim	4
1.5 About the structure of this project	5
1.6 Acknowledgements	5
2 Business Understanding	7
2.1 Business Objectives	7
2.2 Situation Assessment	7
2.3 Data Mining Goals	9
2.4 Project Plan	11
3 Data Understanding	15
3.1 Initial Data Collection	15
3.2 Data Description	16
3.3 Data Exploration	16
3.4 Data Quality	17
4 Data Preparation	19
4.1 Data Selection	19
4.2 Data Cleaning	20
4.3 Data Construction	20
4.4 Data Integration	21
4.5 Data Formatting	21
4.6 Data Updating Process	21
5 Modeling	23
5.1 Modeling Technique	23
5.2 Test Design	25
5.3 Model	25
5.4 Model Assessment	28
6 Evaluation	31
6.1 Model Evaluation	31
6.2 Results Evaluation	33
6.3 Presentation of Results	35
7 Deployment	37
7.1 Current Status	37

7.2	Deployment Plan	38
7.3	Monitoring and Maintenance Plan	38
7.4	Final Report	39
7.5	Project review	39
8	API Reference	43
8.1	src	43
	Python Module Index	57
	Index	59



Note:

This project is an Enexis case on forecasting load demand for distribution transformers.

The results should support grid planners with timely grid reinforcements and replacements.

The case is used as a graduation project for the Professional Education program of the Jheronimus Academy of Data Science

**CHAPTER
ONE**

GENERAL

1.1 About JADS



This project is done in the context of the Professional Education Lead Program of the Jheronimus Academy of Data Science [JADS] as a graduation project.

1.2 About Enexis



Enexis is a regulated regional distribution network operator (DNO) in the Netherlands, responsible for transporting electricity and gas to 2.6 million customers.

Enexis aims to keep the delivery of gas and electricity reliable, affordable and sustainable for all its customers.

1.3 About the current situation

Power is delivered via grid components such as cables and transformers (for electricity) to residential and industrial connections of our customers.

The capacity of these grid components should always be sufficient for the power flows in the grid. This is monitored by Grid Planners and if needed transformers are swapped for heavier ones, or cable connections are strengthened.

For decades this monitoring was done mainly by looking to historical yearly extremes.

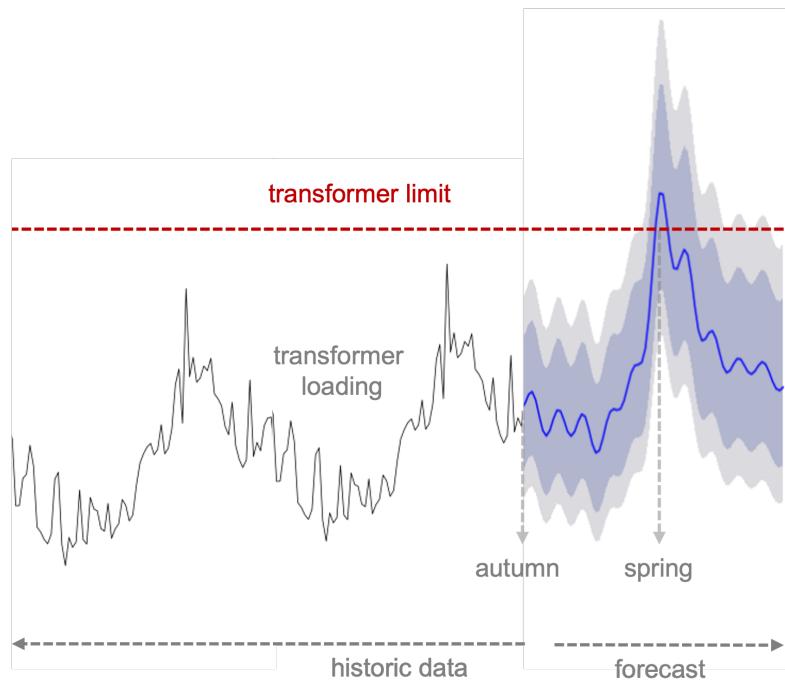
However, the increasing growth rate of power demand and supply driven by emerging technologies as electrical vehicles (EV) and photovoltaics (PV), requires shorter monitoring periods and nowadays even forecasts. Otherwise, there is no time for mitigating actions and customers will be out of power.

1.4 About this project's aim

Enexis started a few years ago with the measurement of its distribution transformer population (in the project "distribution automation light" (DALI)).

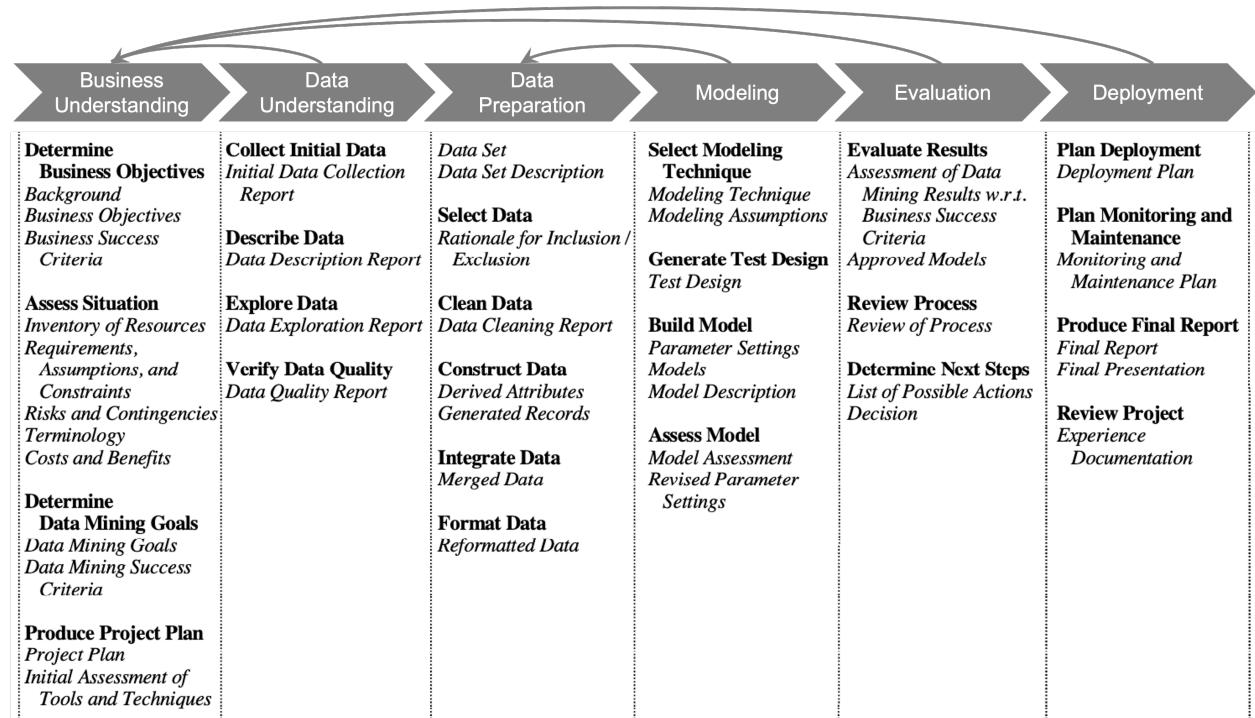
Power measurements are since then available which enable monitoring ad hoc by the grid planners.

This project takes the next step by forecasting on that data in autumn if a transformer overloads in spring.



1.5 About the structure of this project

This project and its documentation are set up around the CRoss Industry Standard Process for Data Mining [CRISP-DM]. It is an iterative process and this documentation focuses on the end result without ignoring the lessons learned along the way. The latter is generally noted at the end of every chapter.



1.6 Acknowledgements

I would like to thank Jeroen the Mast for the valuable feedback and supervision.

Special thanks also goes out to PDEng candidate Akshaya Ravi for her technical support. Together with my buddy David Rijlaarsdam she provided me with helpful insights and discussions on the project.

BUSINESS UNDERSTANDING

2.1 Business Objectives

Enexis is a regulated DNO with a monopoly position regarding the distribution of electricity and gas.

The company wants to keep the delivery of electricity and gas:

- reliable
- affordable
- sustainable

This is realised by doing effective and efficient grid investments according to the Risk and Opportunity Based Asset Management method [ROBAM].

2.2 Situation Assessment

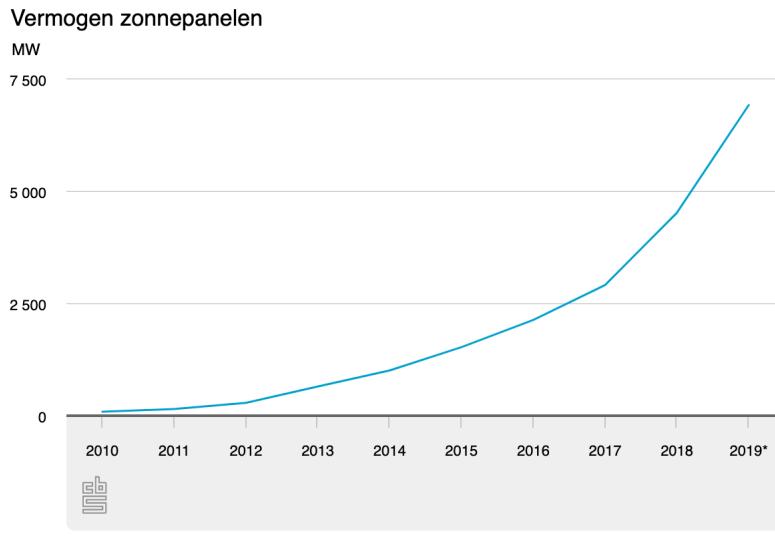
2.2.1 A changing environment

Power is delivered via grid components such as cables and transformers (for electricity) to residential and industrial connections of our customers.

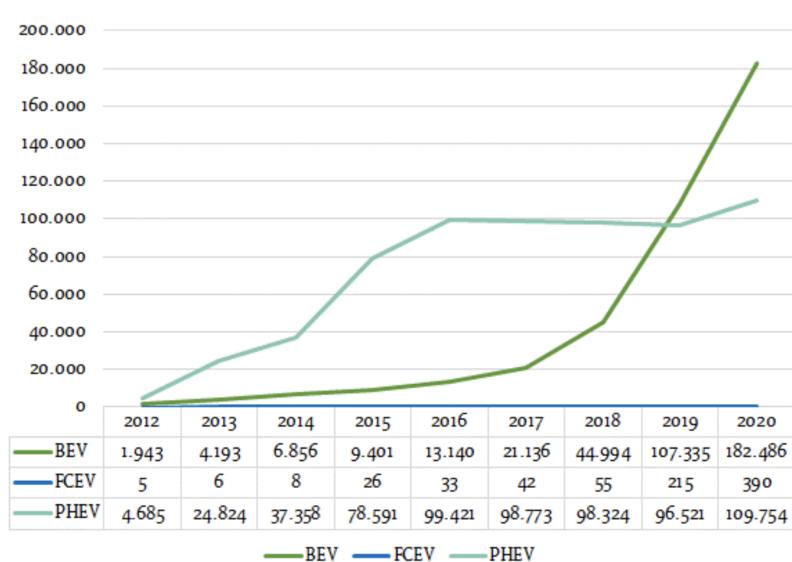
The capacity of these grid components should always be sufficient for the power flows in the grid. This is monitored by Grid Planners and if needed transformers are swapped for heavier ones, or cable connections are strengthened.

For decades this monitoring was done mainly by looking to historical yearly extremes.

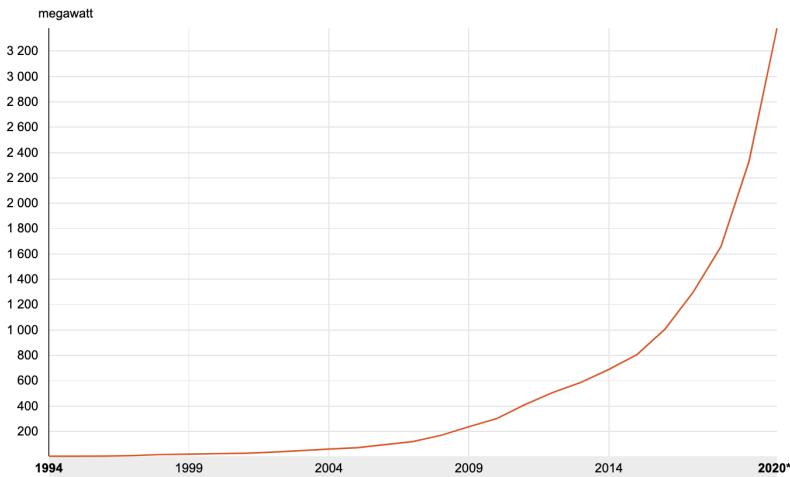
However, the increasing growth rate of power demand and supply driven by emerging technologies as electrical vehicles (EV) and photovoltaics (PV), requires shorter monitoring periods and nowadays even forecasts. Otherwise, there is no time for mitigating actions and customers will be out of power.



Solar power (PV) in The Netherlands (source: CBS).



Battery (B), Fuel Cell (FC) and Plugg-in Hybride (PH) Electric Vehicles (EV) in The Netherlands (source: RVO).



Installed power of heat pumps (HP) in The Netherlands (source: [CBS](#)).

2.2.2 Impact on grid management

The more volatile and more rapidly increasing power flows, require Grid Planners more often to monitor if the minimum and maximum power is still within the transformer's capacity. Otherwise, this will cause unsafe operating situations and it can lead to power outages.

Grid Planners expect a big increase of the number of transformers that reach their operating capacity in the upcoming years. This in a time and environment where Grid Planners are expected to make sure that Enexis is an enabler of the energy transition and not a show stopper.

2.2.3 Opportunity for data science

Since a few years Enexis started to deploy so called Distribution Automation Light Boxes (DALI).

With these measurement apparatus it is possible to monitor our distribution transformers near real time.

At the moment of writing 11k distribution transformers (of the total population of 35k) are equipped with these apparatus which are measuring 15 minute average voltages, currents and powers of each transformer.

This data (together with transformer metadata) enables us to automatically detect or even foresee earlier overloading of transformers. It gives grid planners the opportunity to timely mitigate upcoming issues.

2.3 Data Mining Goals

2.3.1 Primary objective

More volatile power flows require a monitoring tool that forecasts transformer overloading.

The goal is to timely identify future overloading of transformers.

2.3.2 Example Use Case

The project should result in a tool that is able to predict in autumn that a transformer will be overloaded in spring due to EV with a prediction interval.

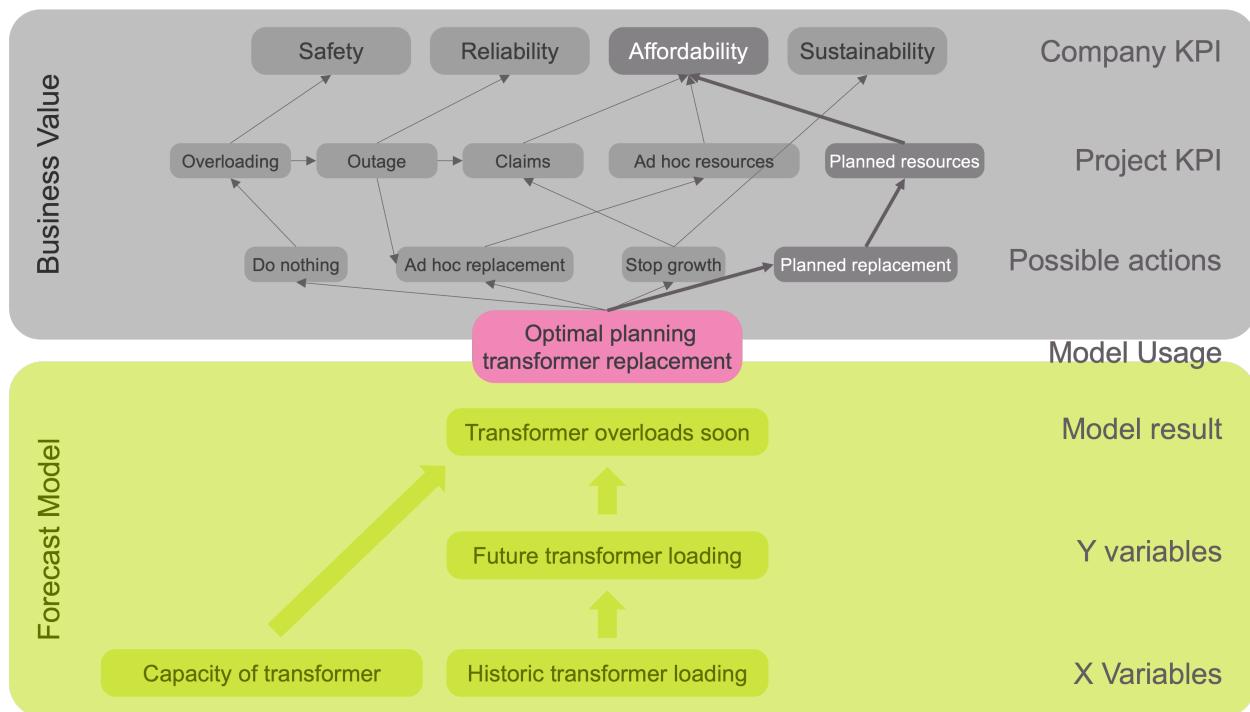
The tool enables better planning of grid strengthening which prevents overloading of transformers (safety and reliability), foreseeing and better plan future work (costs) and enabling the energy transition better (sustainability).

2.3.3 Business Value Diagram

Business value is created by the forecast model by enabling the preferred path for grid investments on the business value part (upper half) of the Business Flow Down Diagram below.

This path is possible by providing grid planners with a probability that a transformer will overload in the foreseeable future.

The model (lower half) will use historic DALI timeseries data to forecast transformer loading. By comparing this forecast with capacity of the transformer possible overloading can be foreseen.



Business Flow Down Diagram for the project.

2.3.4 Requirements from Users

The developed product is only valuable if it is used by the end users, who are the Grid Planners. Therefore the following additional requirements and wishes are important:

- Grid Planners don't want another tool to log into/install..
- The presentation of results has to be quick (no long waiting times).
- Results have to scalable up to 35k transformers.
- The tool should be reliable (stable interface and accurate forecasts).

- The model should be explainable.
- The uncertainty of predictions should be available.
- Results should be available for further use (e.g. importable for load flow tools).
- The results should be sorted based on their urgency.
- Forecast can also be made with limited transformer data.
- Forecast horizon should be six months.

2.3.5 Success Criteria

The project is a success if there is a tool that is being used by the Grid Planners that accurately forecasts overloading of transformers.

Usage of the tool can be measured by tracking users of the tool and by performing interviews with the end users after deployment. Accuracy is assessed by comparing forecasts with measurements.

In more detail the project is a success if:

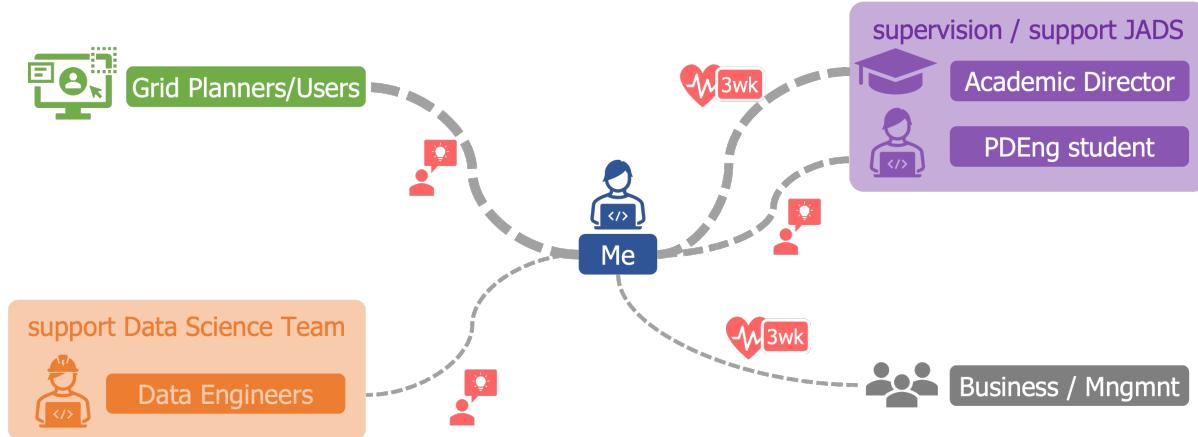
- The model forecasts prediction intervals.
- The working of the model can be explained clearly.
- The prediction intervals ranges are acceptable to the Grid Planners.
- The model can use prior information of the rest of the population if historic measurements are missing.
- The computational burden is acceptable.

2.4 Project Plan

2.4.1 Organisation

The involved stakeholders are:

- Grid Planners: The end users of the tool. Ad hoc, they will be updated/asked for input/their expertise on the field of grid planning and their needs. They assess if the project is a success.
- Data Engineers: To get the tool into production the ICT guidelines and processes within Enexis have to be respected. Expertise of the Data engineers is crucial in the second half of the project to get things into production/deployment.
- Management: The direct manager enables Developer/Lead Bram to work two days a week on this project. Together with all other Enexis stakeholders they will be updated at the end of every sprint in the sprint review session.
- Academic supervision: Jeroen de Mast will be the academic director that monitors progress and the academic level. Every three weeks there will be an one-on-one meeting to discuss progress and issues.
- Academic support: PDEng candidate Akshaya Ravi of JADS is available for support on technical and academic issues. Together with a buddy from the Lead Track periodic meetings (every four weeks) will be planned, but also ad hoc issues will be discussed directly.



The Stakeholders with meeting frequencies (heartbeat / ad hoc).

2.4.2 Resources

The project will take place in June, July, September, October and November 2021.

- **Personnel**

- Project lead / developer: Available 2 days per week.
- Grid Planners: Available 2 hours per week.
- Data engineers: Available 2 hours per week in October - November.
- Supervision / support: Both available for a few hours every 3 weeks.

- **Data**

- DALI 15-minute measurement data.
- Distribution transformer metadata (capacity).

2.4.3 Requirements

The tool should be implemented in such way that it is:

- **Maintainable:** The code has to have docstrings and unit tests.
- **Scalable:** The tool has to be able to process 35k transformers.
- **Deployable:** The tool has to be able to go in production according to the Enexis standards (Test-Acceptance-Production).

2.4.4 Assumptions

- DALI data is available without huge quality issues during the project.
- Weekly extremes on transformer data is an acceptable aggregation level for capacity planning.
- Data Engineering has capacity for several hours a week for support between September and December.
- The computational burden for probabilistic models is feasible regarding the computational power available.

2.4.5 Risks and Contingencies

- Lead / developer has just become a father (is technically up to September on parental leave) and bought a house that has to be decorated.
– Mitigation: None.
- Grid Planners are immensely preoccupied with the current challenges in the grid. Although not a lot of time is required, it is essential to be clear and direct regarding expectations and communications and limit the effort and time for this project for Grid Planners without giving in on quality/input.
– Mitigation: Be clear and direct regarding expectations and communications and limit the effort and time for this project for Grid Planners without giving in on quality/input.
- Data Engineers are also loaded with work and might not have time/resources available.
– Mitigation: It is essential to request capacity in the beginning of the project, although they will be involved only in the second half.
- There is no use of sensitive data in this project regarding privacy (GDPR) or security. DALI data is allowed to be used. Credentials are not embedded in code and access to data sources is restricted by design.

2.4.6 Costs and Benefits

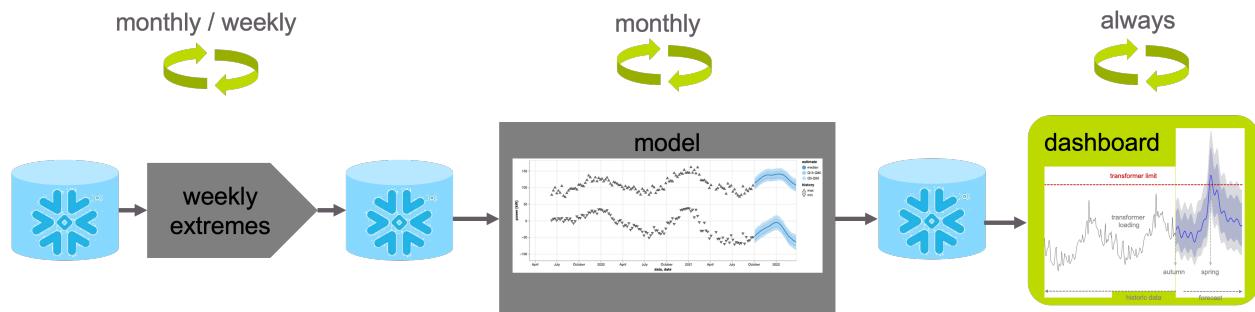
- Data collection: Data is available in an existing database. Only querying is needed, no active additional data collection.
- Implementation: Open source software is used besides already licensed applications. Only computation power will cost additionally. More details will be available after a first proof of concept.

2.4.7 Implementation Concept

The tool is split into three steps to tackle the amount of data available, but still make the results manageable:

- Preprocessing: To condense 15 minute load averages into weekly extremes (minimum and maximum).
- Forecasting: To fit a model on the aggregated data and create forecasts.
- Dashboarding: To display (forecast) results to the end user.

In between steps results are stored in a Snowflake database.



The proposed steps for implementation of the tool.

2.4.8 Planned Milestones

Table 1: Scheduled project milestones for 2021.

week	CRISP-DM step	detail
26	Business Understanding	
29	Data understanding	Go / No Go
-	-	Summer break
36	Data Preparation	
39	Modeling	
42	Evaluation	Go / No Go
45	Deployment	

DATA UNDERSTANDING

3.1 Initial Data Collection

The goal is to forecast accurately maximum and minimum loading of transformers.

There is 15-minute average data available of 11k transformers.

The maximum and minimum loading is not needed on this detailed level. According to the Grid Planners, weekly or monthly extreme details are more than enough.

For this first model to forecast transformer loading no external variables will be used as input for the model. Only historic data will be used (with metadata to extract de capacity of the transformer).

Data is available in Snowflake. A Jupyter Notebook pandas with a SQLAlchemy engine is used for the data collection from Snowflake.

Credentials are available via Vault in the Enexis domain.

3.1.1 Measurement Data Source

In a Snowflake database the following data is available:

Table 1: Snowflake data source details.

Database	DB_CDWH_P
Schema	CDWH_4_BDM
Table	BDM_DALI_METINGEN

3.1.2 Metadata Data Source

In a Snowflake database the following metadata is available:

Table 2: Snowflake data source details.

Database	DB_CDWH_P
Schema	CDWH_5_INF_MEETDATA_E DIM_DALI_BOX_VW
Table	DIM_DALI_BOX_VW

3.2 Data Description

3.2.1 Measurement Data Details

The measurements are available since Q2 2018.

Since then increasingly more transformers were equipped with DALI boxes and at the moment of writing (Q3 2021) a total number of 10,992 boxes are measured. Measurements on open doors (from safety perspective important) and currents, voltages and powers accumulated up to 89,052,020,404 records in the Snowflake table.

For this project we focus first on the active power on a transformer on the medium voltage connection side.

The power is available for all three power phases as well as the sum of them.

Below are the fields of interest given that are queried from the aforementioned table.

Table 3: DALI table fields of interest.

Field o.i.	Type	Example
BOXID	VARCHAR	ESD.000240-2
CHANNELID	VARCHAR	register://electricity/0/activepower/sumli?avg=15
WAARDE	DOUBLE	-3.408062
DATUMTIJD	TIMESTAMPTZ	2021-05-12 07:45:00.000000000

3.2.2 Metadata Details

In the metadata table there are 15,058 records present.

Only 9,920 boxids are operational and have a nominal power correctly registered.

Table 4: Metadata table fields of interest.

Field o.i.	Type	Example
BOXID	VARCHAR	ESD.000240-2
BOX_BEDRIJFSSTA-TUS	VARCHAR	in bedrijf
IN_BEDRIJFS-NAME_DATUMTIJD	TIMESTAMPTZ	2021-05-12 07:45:00.000000000
VESTIGING	VARCHAR	Breda
VERMOGEN_NOMINAAL	NUMBER	400

3.3 Data Exploration

The measurement data is up-to-date, but has missing values in two moments in time (both in 2021). This issue was discovered and is addressed.

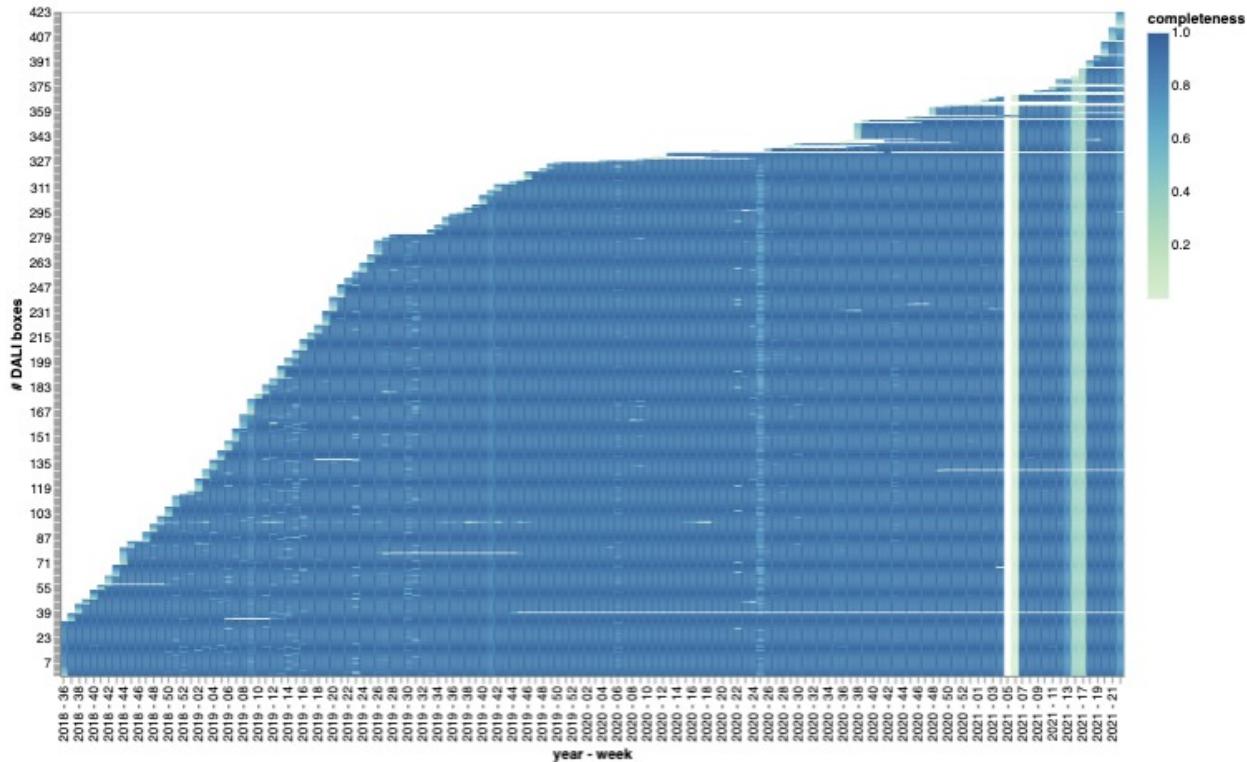
For the project it will not be a huge issue, since the data is aggregated to weekly minimum and maximum values. It will cause these extremes to be less severe if (a lot of) data is missing. Only if data is missing for a whole week, there will also be missing data in the aggregated set.

In both cases (outlier and missing data) the model can handle this.

For time series modelling it is advised to have at least two periods (years in this case) of measurement data.

As the figure below shows, this might be a problem eventually. One of the challenges is therefore to explore if prior knowledge of the population can overcome this issue.

In the figure below the completeness of the data over time is given for the transformers in the service area of Breda. On the vertical axis the transformers are shown ordered by the time they became operational. On the horizontal axis the time is given. In color the completeness is given per week: Darkblue (value 1) means that all data was present, towards white (value 0) means no data at all.

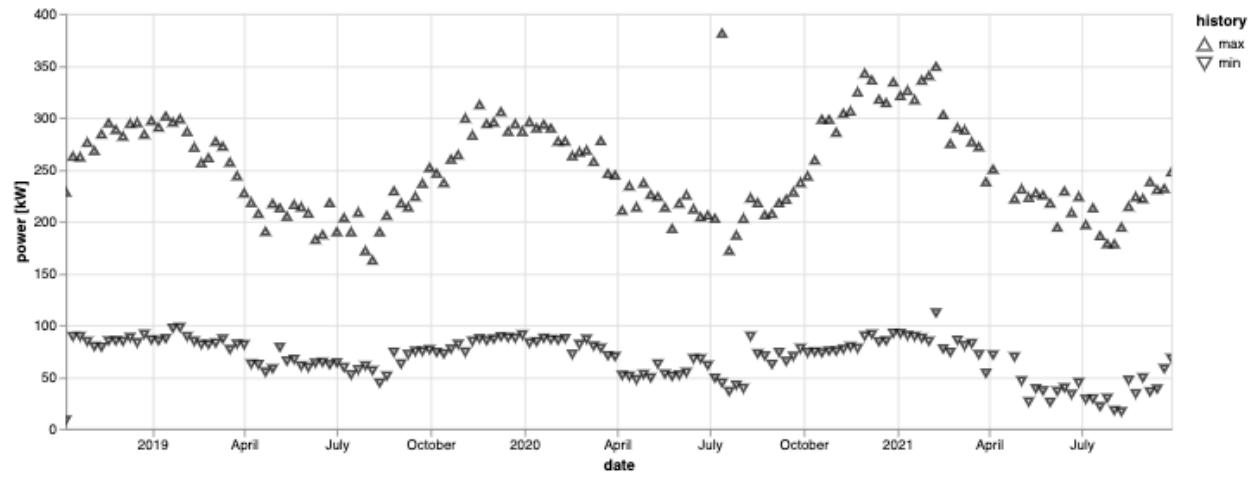


Completeness for DALI data in the service area of Breda.

3.4 Data Quality

Beside the missing data described above, the data quality (of the 15-minute power averages) is as like expected. The reason is likely the 15-minute averaging already smooths out the extreme (short circuit) values and measurement errors.

Although sometimes outliers can still be seen in the data (which can propagate into the weekly extremes as shown in the figure below). Taking not only the extremes, but also the second highest/lowest value per week for robustness did not make a lot of difference (probably also due to the aforementioned smoothing).

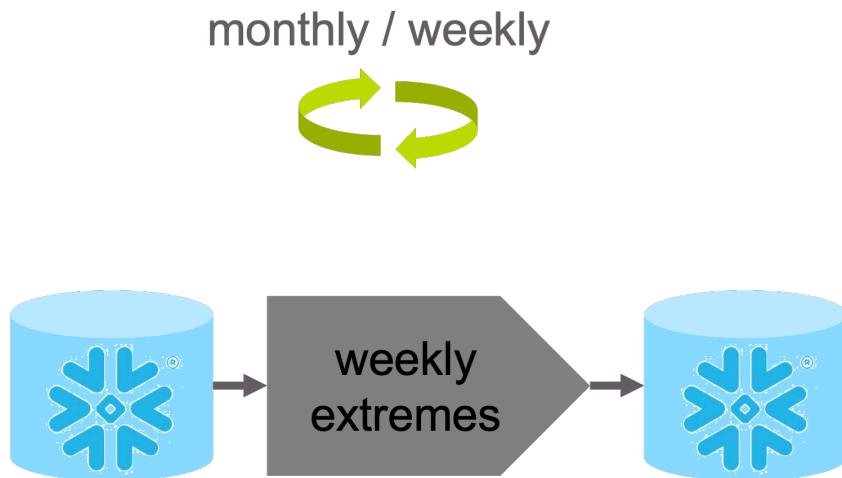


Example of weekly extremes with an outlier for the maximum in July 2020.

DATA PREPARATION

The data preparation step focuses on converting raw measurement data with a frequency of 15 minutes into weekly extremes and preparing these extremes to be usable input for the forecasting model.

Initially the entire history is aggregated and stored in a Snowflake database. Consecutively, updates are done every week.



The data preparation step resulting in weekly aggregated data.

4.1 Data Selection

Columns described in Data Understanding are selected from the measurement data and metadata tables.

From the measurement data only data is selected from DALI boxes that have nominal power registered in the metadata table and that are in operation (`src.utils.snowflake.read_meta()`). Both data and metadata are needed to indicate future overloading and add value to Grid Planners.

As mentioned only the active power (P) on medium voltage side is used initially. Apparent power (S) is preferred (for a more fair comparison with the nominal power), but since this is 15 minute average data, this is hard to reconstruct.

In this step the separate power phases are selected and processed as well as the sum of the phases (`src.utils.snowflake.make_week_extremes_query()`).

Only the 15 minute average channels are selected for preprocessing.

4.2 Data Cleaning

No data cleaning is performed on the raw data before aggregation, but data is checked and cleaned on the following after loading (`src.preprocess.preprocess.load_data()`) and before being used by a model for forecasting:

- Data of extremes (minimum or maximum) having the value of zero in the beginning of the series are removed (`src.preprocess.preprocess.remove_leading_idling()`). This is for example the case if a DALI box is in operation, but its transformer is not.
- Only data is used that has a history of more than two years (`src.preprocess.preprocess.too_short()`). This will ensure in this stage that the seasonality (sub)model has enough data to tune on.
- Only data is used of transformers that have a measurement in their history with an absolute value higher than half the transformer capacity (`src.preprocess.preprocess.too_small()`).

Duplicate data (can only be created by updating the extreme table) is not an issue for the model and will not be eliminated. Missing data is neither a problem for the model and is also not imputed.

4.3 Data Construction

From the raw 15-minute data the weekly minimum and maximum are determined. This is done per channel and boxid (`src.utils.snowflake.make_week_extremes_query()`). The week definition used is the ISO-week since this is always a full week.

A SQL query aggregates and writes the result asynchronously on the Snowflake database. This can be done in batch for all historic measurements (`src.utils.snowflake.create_week_extremes()`), but the created table can also be updated per week (`src.utils.snowflake.update_week_extremes()`).

Table 1: Snowflake table details for weekly extremes data.

Database	DB_DATASCIENCE_P
Schema	DATASCIENCE_1_ETL
Table	DS_SPARK_DALI_WEEK_EXTREMES

The fields of the table are listed below. The table is clustered by BOXID and L (phase). The amount of rows is condensed from 89,052,020,404 to 3,457,856 records.

Table 2: Extremes table fields.

Field	Type	Example
BOXID	VARCHAR	ESD.000240-2
L	VARCHAR	sumli
YEAR	NUMBER	2021
WEEK	NUMBER	53
PROCESSED_ON	TIMESTAMPTZ	2021-05-12 07:45:00.000000000
MAX	DOUBLE	678.90
MIN	DOUBLE	123.45

4.4 Data Integration

Since no additional data sources are used, no joins or merges are required.

4.5 Data Formatting

The model does not demand an order (e.g. by year and week) of the data. For the modelling stage the data is queried from the table in *Data Construction*

Consecutively, a date column is constructed from the ISO year and week format with day==1.

The extra columns period and model_var are assigned and filled with the values “history”, “observed” respectively for measurement data. This is in preparation for long formatting and concatenating forecast results in a later stage (`src.preprocessing.format_data()`).

An example of the loaded extreme data is shown below:

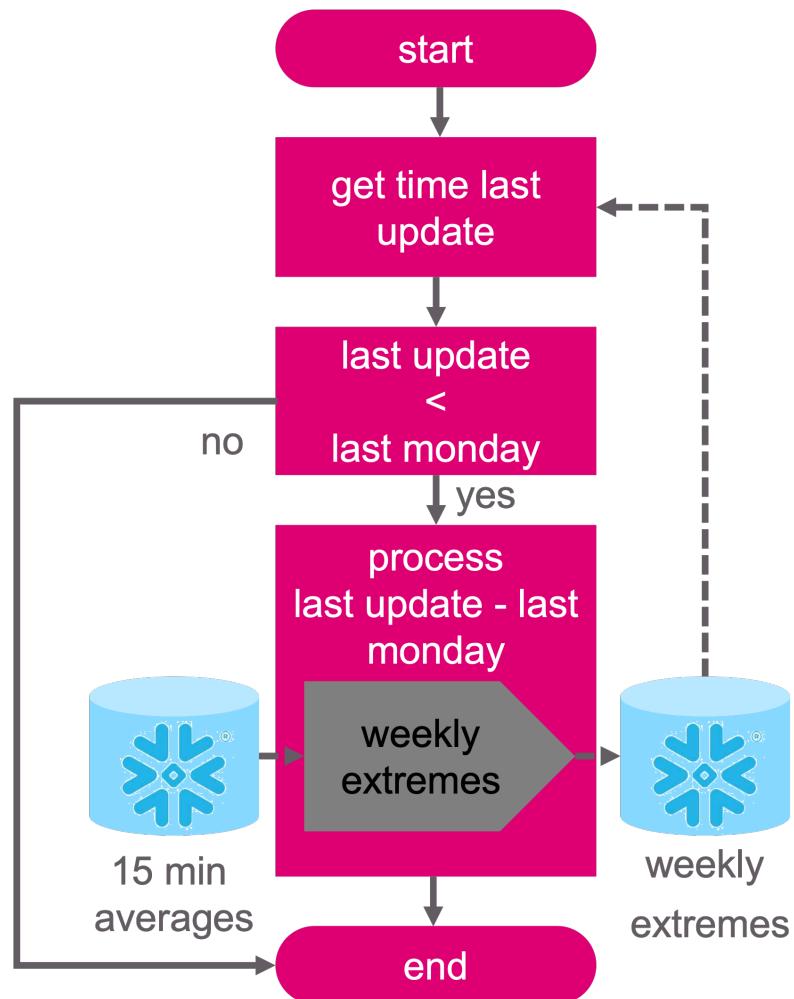
df_data										
boxid	date	I	processed_on	week	year	extreme	value	period	model_var	
0	VRY.CHOPS-1	2018-09-03	sumli	2021-10-07 07:03:50.602000+00:00	36	2018	max	145.486206	history	observed
1	VRY.CHOPS-1	2018-09-10	sumli	2021-10-07 07:03:50.602000+00:00	37	2018	max	147.918304	history	observed
2	VRY.CHOPS-1	2018-09-17	sumli	2021-10-07 07:03:50.602000+00:00	38	2018	max	154.419098	history	observed
3	VRY.CHOPS-1	2018-09-24	sumli	2021-10-07 07:03:50.602000+00:00	39	2018	max	155.272598	history	observed
4	VRY.CHOPS-1	2018-10-01	sumli	2021-10-07 07:03:50.602000+00:00	40	2018	max	169.614395	history	observed
...	
313	VRY.CHOPS-1	2021-08-30	sumli	2021-10-07 07:03:50.602000+00:00	35	2021	min	-223.345993	history	observed
314	VRY.CHOPS-1	2021-09-06	sumli	2021-10-07 07:03:50.602000+00:00	36	2021	min	-211.977295	history	observed
315	VRY.CHOPS-1	2021-09-13	sumli	2021-10-07 07:03:50.602000+00:00	37	2021	min	-209.037598	history	observed
316	VRY.CHOPS-1	2021-09-20	sumli	2021-10-07 07:03:50.602000+00:00	38	2021	min	-223.494995	history	observed
317	VRY.CHOPS-1	2021-09-27	sumli	2021-10-07 07:03:50.602000+00:00	39	2021	min	-173.239502	history	observed

The format of the loaded extremes data.

4.6 Data Updating Process

The weekly extremes can be updated on a weekly (or longer) basis.

By running `src.preprocessing.update_extremes()` the function `src.utils.snowflake.update_week_extremes()` is called. This will trigger the following steps which update the weekly extremes Snowflake table:

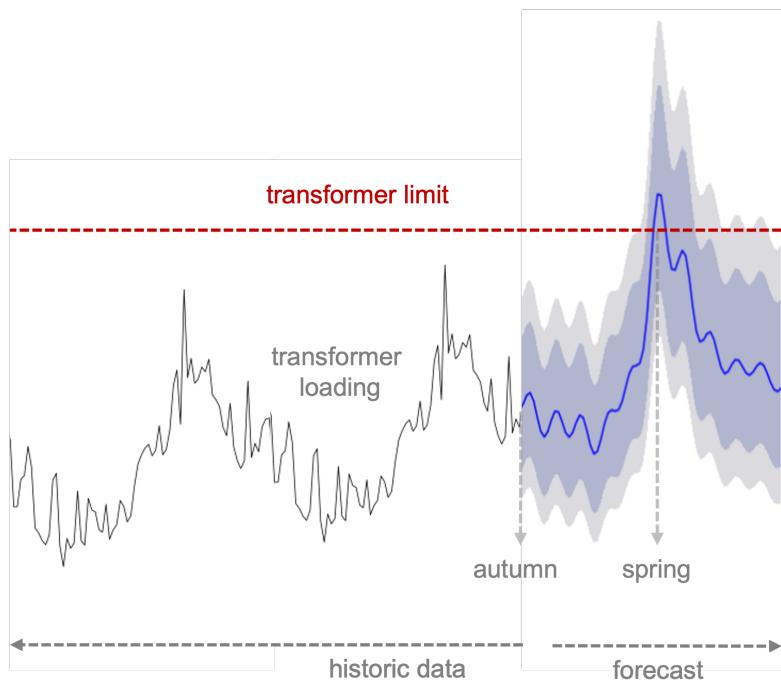


The detailed process to create and assess load forecasts.

MODELING

5.1 Modeling Technique

The Grid Planners are looking for a data driven forecast with a uncertainty indication as mentioned before. This is visualized as a concept in the figure below.



A possible output of the model visualized.

5.1.1 Input Features

The Grid Planners want to have a result dat is based on real observed data. Or, to make it more explicit: Data that has been measured and not data from possible scenarios with a wide range of uncertainty.

Another suggestion was made to use (historic) weather data as an exogenous feature / input for the model. However, this is not data we have available with a reasonable confidence for the forecast horizon (in the future). And seasonality can also be captured from the observed data itself for different timeseries models.

Therefore, the model only uses historic DALI measurement data.

5.1.2 Fit-Predict Process

A model will be fitted / tuned for every DALI measurement extremes since every transformer has it's own unique signal (`src.forecast.forecast.determine_estimates()`).

A model is trained on observed data, a forecast is made and stored and finally the trained model is discarded (`src.forecast.forecast.forecast()`).



The weekly extremes are loaded for tuning the model. The model forecasts are afterwards stored in a Snowflake database. The next time a forecast is required, new data is available and a new model will be fitted and used for forecasting.

5.1.3 Model Environment

A probabilistic approach has been implemented to fulfill the wish to forecast and display uncertainty rather than a point estimate.

From the main probabilistic toolboxes (PyMC3, STAN, EDWARD2, Pyro, TensorFlow2 Probability) PyMC3 was used for it's extensive documentation.

A future step could be to transfer the model into TensorFlow2 Probability since PyMC3 is not the most recent toolbox anymore.

5.1.4 Generalized Additive Model Concept

To address the explainability of the model, a similar approach as Facebook Prophet is used.

The model is a Generalized Additive Model (GAM) that consists of a trend / drift and seasonality component (and an error component).

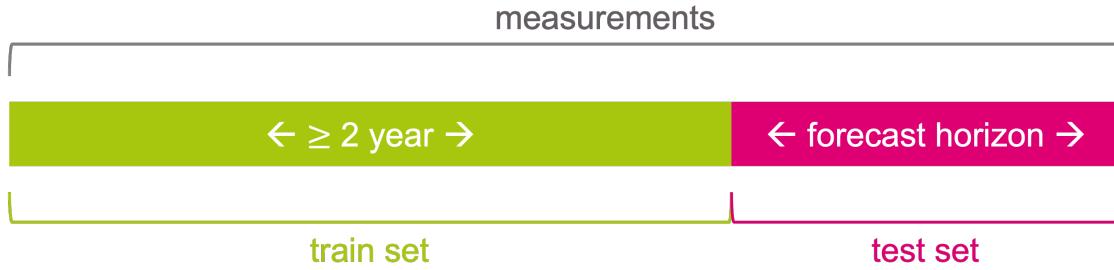
The translation to PyMC3 of Richie Vink was used as a basis.

The GAM approach makes it easy to decompose the different components of the timeseries and show the Grid Planners the effect of the drift and seasons separately.

5.2 Test Design

To evaluate the model the observed data is split into a train and test set based in the forecast horizon (`src.preprocess.preprocess.split_last()`).

After splitting the train set is tested again for not being too short (at least two years: `src.preprocess.preprocess.too_short()`).



The split of measurement data into train and test data.

The test set is only used to validate forecasting results.

The train set is used to train / fit / tune the model.

5.3 Model

5.3.1 Generalized Additive Model

The GAM model used is (`src.model.model.create_model()`):

$$\sigma_\epsilon \sim Uniform(lower = 0, upper = 1)$$

$$\Sigma | drift, yearly, \sigma_\epsilon = Normal(\mu = drift + yearly, sd = \sigma_\epsilon)$$

The additive naming is explicit in this notation.

The error component has a has a bandwidth of σ_ϵ .

Drift Component

According to the Grid Planners a increasing growth is more and more common due to the energy transition.

Therefore, a stable drift model (`src.model.model.drift_model()`) is used that can mimic that.

An exponential function resulted in divergence during the model tuning, but a second order taylor series makes the model convergent and stable.

The drift component model with a taylor series with the order of n is described as:

$$X_{drift}(t) = [t^0, \dots, t^n]$$

$$\beta_{drift} \sim Normal(\mu = 0, sd = 0.5)$$

$$drift | \beta_{drift} = X_{drift}(t) \beta_{drift}$$

For modelling a drift that has the described growth, a polynomial with order $n = 2$ is used.

Yearly Component

Since the data has been aggregated into weekly extremes, the only seasonality to model is the yearly pattern.

The yearly seasonality is modeled with n order fourier series (`src.model.model.seasonality_model()`).

This is based on the work of [Richie Vink](#).

The yearly seasonality model is described as:

$$X_{yearly}(t) = [\cos(\frac{2\pi 1t}{T}), \dots, \sin(\frac{2\pi nt}{T})]$$

$$\beta_{yearly} \sim Normal(\mu = 0, sd = 1)$$

$$drift | \beta_{yearly} = X_{yearly}(t) \beta_{yearly}$$

Here the T is the period of the seasonality in unit of time of the data.

The unit of time in this case is a week for the data and a year in weeks is $T = 52.1775$.

The order taken for the fourier series is $n = 5$.

Enabling Forecasts

The model parameters (β)'s can now be tuned to produce the most model to produce the observed (measurement) data.

To forecast, the model also needs to produce samples beyond the timestamps it has been tuned on.

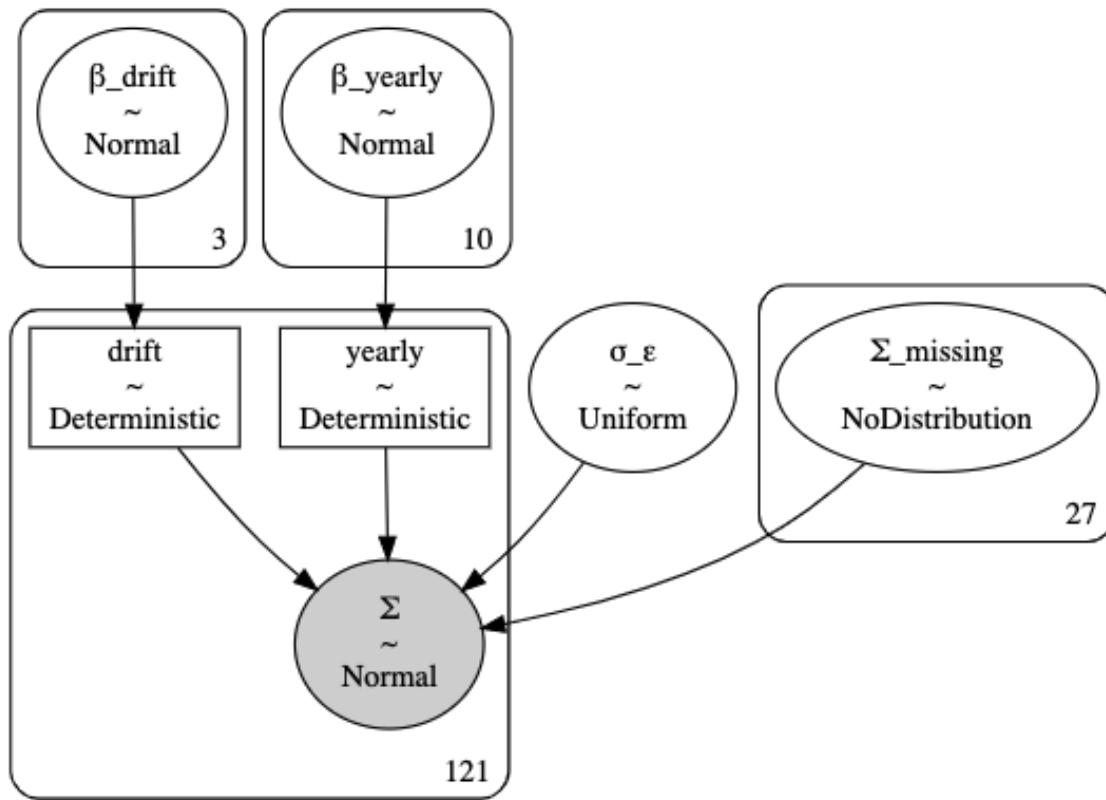
The PyMC3 model can cope with this by feeding it with timestamps that are extrapolated for the forecasting horizon (`src.preprocess.preprocess.extrapolate_timestamps()`).

The matching observations (measurements) can be intentionally filled with NaN's.

In the model PyMC3 will name them $\Sigma_{missing}$. This characteristic makes the model also robust against missing data.

By sampling the posterior predictive after tuning, also samples are generated for the extrapolated forecast timestamps (`src.forecast.forecast.determine_estimates()`).

Total model Σ



The total model visualized.

Two separate GAM models Σ (`src.model.model.create_model()`) are used for the weekly minimum and maximum.

The visual above shows the total GAM model with a polynomial drift order $n = 2$ (the bias of order 0 explains $n + 1 = 3$) and a fourier order of $n = 5$ (the sine and cosine parts explain $n * 2 = 10$).

The number of observations (weeks of measurements for this case) is 121 and the forecasting horizon is just more than six months (27 weeks).

Formatting Results

From the model posterior predictive samples are drawn for all timestamps (also measurement timestamps, 1000 samples per timestamp).

From the posterior samples, the quantile bands are determined (`src.forecast.format.make_quantile_bands()`). This reduces the data storage.

The upper and lower limits of the quantile bands are then stored in the same format as the input (`src.forecast.format.format_model_estimates()`).

The input of the model and the output are then concatenated together. This eases the visualization later.

df_total													
	boxid	date	I	processed_on	week	year	extreme	value	period	model_var	band	boundary	
0	HVT.1111153-1	2019-06-10	sumli	2021-10-07 07:03:50.602000+00:00	24	2019	max	78.797707	history	observed	NaN	NaN	
1	HVT.1111153-1	2019-06-17	sumli	2021-10-07 07:03:50.602000+00:00	25	2019	max	91.941017	history	observed	NaN	NaN	
2	HVT.1111153-1	2019-06-24	sumli	2021-10-07 07:03:50.602000+00:00	26	2019	max	100.725197	history	observed	NaN	NaN	
3	HVT.1111153-1	2019-07-01	sumli	2021-10-07 07:03:50.602000+00:00	27	2019	max	83.487396	history	observed	NaN	NaN	
4	HVT.1111153-1	2019-07-08	sumli	2021-10-07 07:03:50.602000+00:00	28	2019	max	91.487610	history	observed	NaN	NaN	
...	
721	HVT.1111153-1	2021-08-30	sumli		NaT	35	2021	max	129.740717	future	Σ	median	lower
722	HVT.1111153-1	2021-09-06	sumli		NaT	36	2021	max	133.958580	future	Σ	median	lower
723	HVT.1111153-1	2021-09-13	sumli		NaT	37	2021	max	137.037553	future	Σ	median	lower
724	HVT.1111153-1	2021-09-20	sumli		NaT	38	2021	max	140.406872	future	Σ	median	lower
725	HVT.1111153-1	2021-09-27	sumli		NaT	39	2021	max	145.338946	future	Σ	median	lower

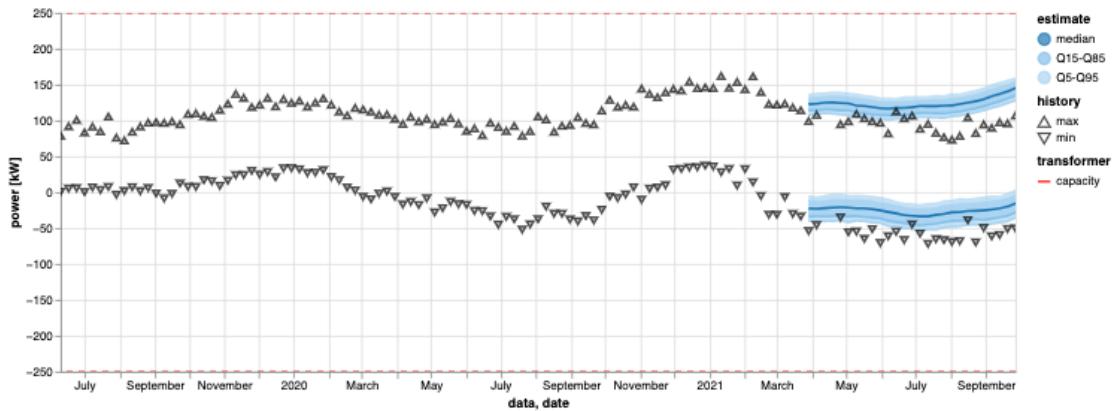
The concatenated input and result.

5.4 Model Assessment

The following model findings are most salient:

- **The model converges during tuning and gives feasible results.**
 - Exponential drift function tuning will not converge.
- **The computational burden on a CPU to tune and forecast both extremes is 1:24.**
 - CPU: 2 GHz Quad-Core Intel Core i5
 - RAM: 16 GB
- The model is fairly insensitive to outliers and missing data.
- The splitting of observations into train and test set works.
- The extrapolation with the forecasting horizon works.
- **An pure additive model may not be sufficient.**
 - Growth also increases the yearly component (see visualization below).
 - A pure multiplicative diverges.
 - A hybrid model (addition of a small fraction of a multiplicative model) might be an option.

A visualization of the results is shown in the figure below which shows most of the aforementioned points:



An visualization of the measurements (history) and forecast (estimates). Measurements from the train and test set are plotted.

5.4.1 Improvement suggestions

The following ideas could result in a better model:

- Implementing a hybrid additive-multiplicative model for dealing with the growing seasonality.
- **Adding a extra component to detect temporarily bypass switching of loads of other transformers.**
 - This could be implemented by estimating parameters of a rectangular function.
- **Making more recent observations more relevant for slowly changing loading patterns.**
 - Possibilities are to mimic weights with `pm.Potential` or `pm<distribution>(tau=weights)`.
- Using the population seasonality as a `prior` in case of a short history of observations.
- Using by-pass dummy model for `outlier` robustness.

**CHAPTER
SIX**

EVALUATION

This section discusses the evaluation of the implemented model, the assessments of results and a reflection on the learning objectives for the study course this project was part of.

6.1 Model Evaluation

The evaluation of results has not been done with a error criteria (e.g. MAE, MAPE), but only visual. This will be done in the near future, when additional model improvements are implemented and needed to be assessed.

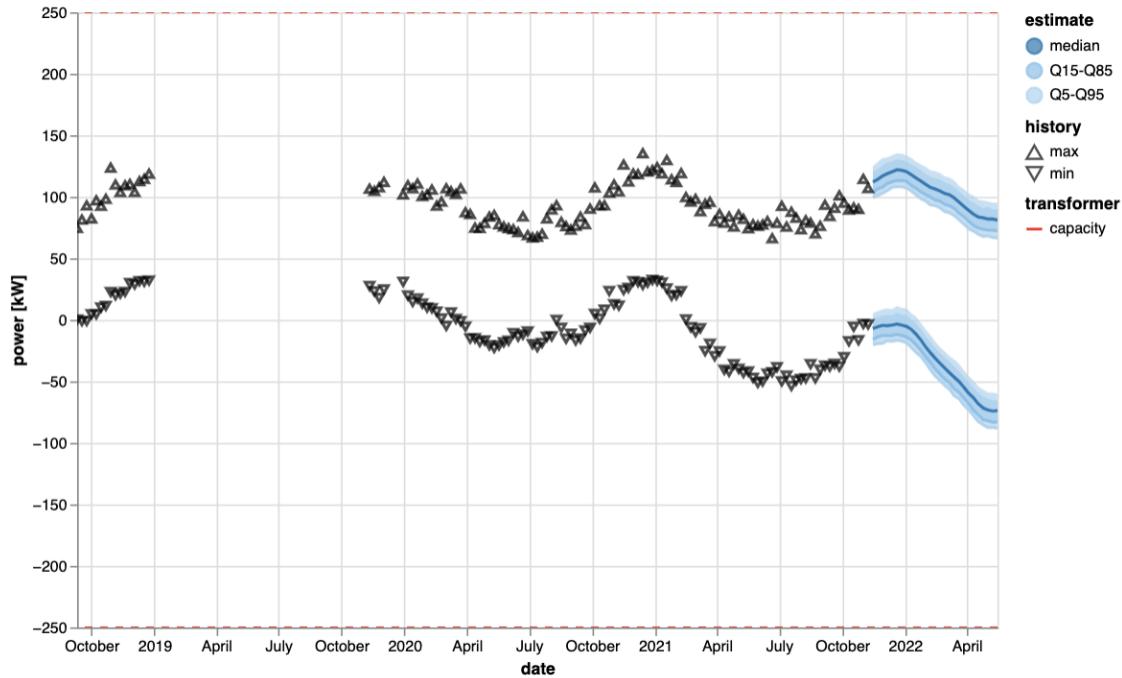
This section will discuss briefly the model

6.1.1 Model Robustness

The model as it is described before is stable for all transformer forecasts: It converges for all transformer timeseries that are used.

A robust model for all 11k transformers is essential if models are trained and used without human supervision and results are automatically assessed. An experiment of an exponential function for the drift sub-model was implemented, but far from stable.

The model is also robust when data is missing. Even missing data for long periods will result in a reasonable forecast:



The model results after data accidentally was removed from the source database.

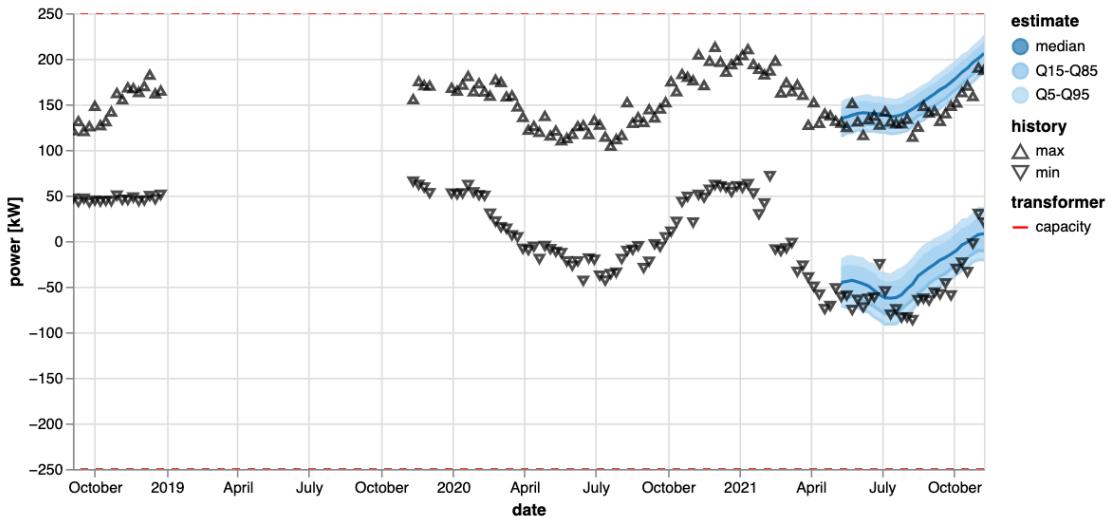
6.1.2 Model Results

The model forecasts neatly the median and quantile bands as intended.

However, it will always be a model trained on historical, observed data, and forecasts will always be off. Partially, this can be due to model imperfections, for which improvements are suggested earlier, partially this is due to the reality of the outside world.

The grid planners were shown different forecast that deviated from realisations to make them aware of the limiting accuracy and reliability of the models and their forecasts.

This was visually shown by the following actions: The results were assessed by splitting the observed data into a train and test set, using a model based on the train set and comparing the results with the test set. The performance for the different transformers of course also differs, but an example can be seen in the following figure:



Forecasts made by a model based on data before May 2021, validated with later observation.

6.1.3 Model Improvements

Stepping through the CRISP-DM cycle resulted in several insights to for model improvement. The most promising suggestions for model improvement are:

- Improving tuning and forecasting time.
- Implementing hybrid additive-multiplicative model for dealing with the growing seasonality.
- Adding a extra component to detect temporarily bypass switching of loads of other transformers.
- Making more recent observations more relevant for slowly changing loading patterns.
- Using the population seasonality as a *prior* in case of a short history of observations.
- Using by-pass dummy model for *outlier* robustness.

6.2 Results Evaluation

This section discusses the assessment and prioritization of the forecasts and the implementation within the process.

6.2.1 Results Prioritization

There are potentially 11k models that all forecast six months ahead.

These are numbers too high to be assessed by grid planners one by one. Therefore the forecast results are automatically assessed and ordered by urgency.

The grid planners gave as input that they wanted overloaded transformers ordered by the time the potential overloading was expected.

Firstly, the definition of potential (over)loading was agreed on to be the following:

- Potential loading is the maximum absolute of the forecast quantile bands divided by the transformer capacity.
- Potential overloading is the potential loading is greater than one.

Potential overloaded transformers can now be ordered by the point in time when they reach this limit.

Then there are also transformers that will never potentially overload. The forecasts of these transformers are simply ordered by potential loading.

To summarize into steps:

- Determine absolute value of the forecasted quantile bands.
- Divide the found value by the transformer capacity.
- Clip the results on the value of one.
- Find value and index (timestamp) of the maximum.
- Order the transformers by timestamp (first) and value (second).

The result of this prioritization is shown below.

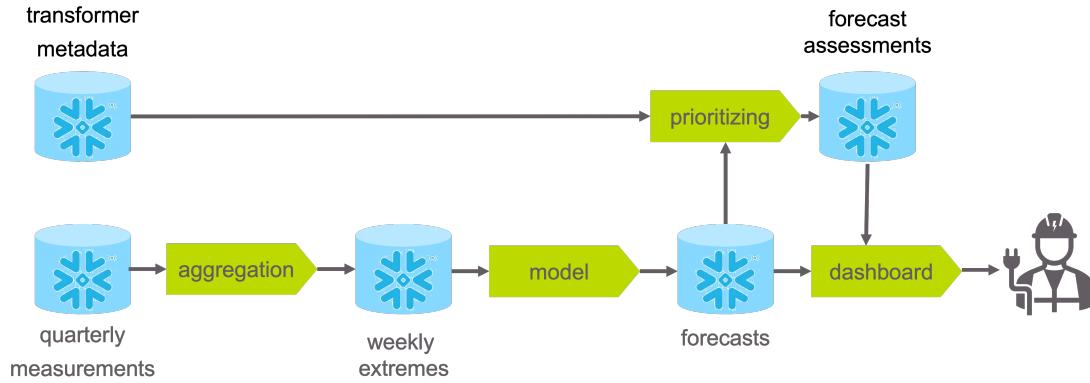
	boxid	P_max	t_max
0	VRY.HOUBE-1	1.000000	2021-11-15
1	BRAM.VONK-1	1.000000	2021-11-19
2	063.691-1	0.875502	2021-12-13
3	159.628-1	0.827228	2021-12-13
4	ESD.000363-1	0.778692	2022-02-14
5	VRY.MHANS-1	0.756287	2021-12-13
6	ESD.000042-1	0.679398	2022-02-07
7	154.732-1	0.678712	2022-02-07
8	121.633-1	0.667767	2021-12-27
9	121.670-1	0.641002	2021-11-22
10	133.235-1	0.580517	2022-01-24
11	ESD.000451-1	0.577681	2022-02-28
12	063.555-1	0.565492	2022-01-17
13	063.669-1	0.541530	2021-12-20
14	VRY.SCHOS-2	0.539767	2021-11-15
15	ESD.000090-1	0.517359	2022-01-17
16	ESD.000319-1	0.477418	2022-01-31
17	154.730-1	0.430946	2021-11-22
18	VRY.HANDR-1	0.358135	2022-01-24

Transformers ordered by potential (over)loading.

The result is that grid planners have an prioritized list of transformers, with on top the transformers that are probably overloading soon.

6.2.2 Evaluation Implementation

The assessments of the forecasts which results in the ordered list (`src.forecast.assess.asses_forecasts()`) is done directly after forecasting.



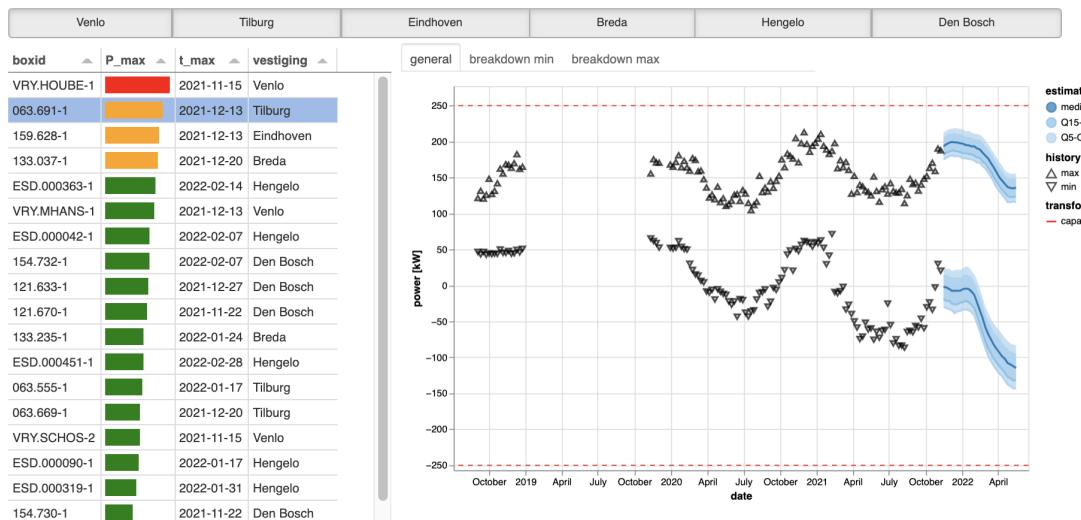
The process steps of the final product.

The result is stored separately in a Snowflake database table (`src.utils.snowflake.read_forecast_meta()`). This way the data can quickly be loaded by the dashboard application and presented to the grid planners.

6.3 Presentation of Results

A dashboard is made by using Panel and Altair libraries (`src.plot.dashboard`).

The dashboard presents the ordered list of all transformers and a figure of a forecast for an individual transformer.



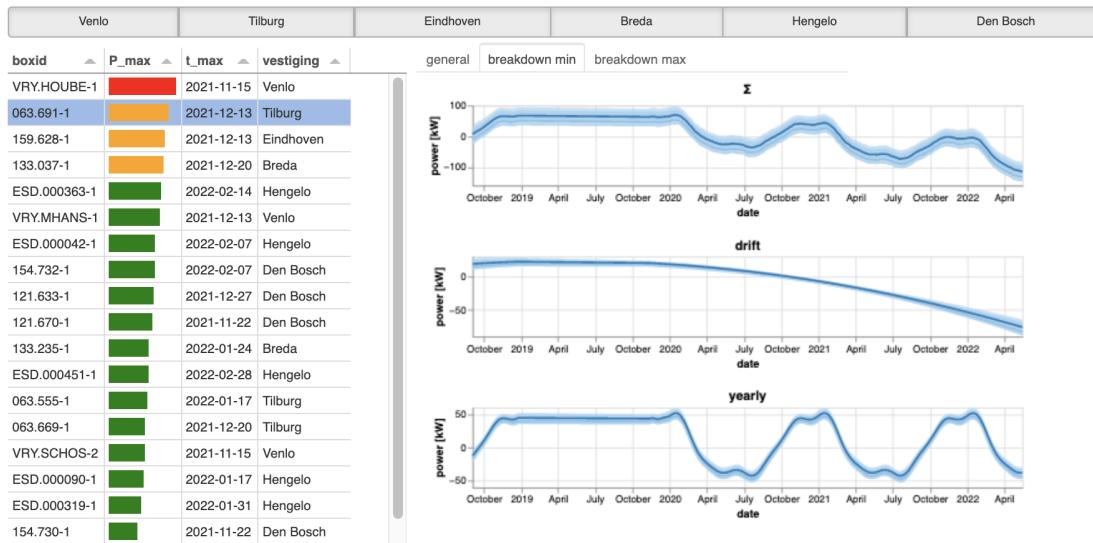
The dashboard with on the left the ordered transformers and on the right the tabs with forecast and decomposed trend and yearly pattern for minimum and maximum.

The list can be filtered by Enexis sub-service area, since grid planners are generally responsible for an area within these areas and not interested in the whole service area. Other ordering is possible if the grid planner desires.

After selection of a transformer, the measured weekly extremes, the forecast six months ahead, and the transformer capacity are displayed in the first tab.

One can zoom in and pan de plot and tooltips show up on historic measurement data.

If the grid planner is interested, he can view in the other tabs the decompositions (drift / trend and yearly pattern) for the weekly minimum and maximum.



The dashboard with the sub-service area filtering and the minimum decomposition tab of the dashboard.

6.3.1 Improvements on Presentation

Firstly, the loading of individual forecasts requires some improvement. At this moment the data loading is slow due to an interface issue between python and the Snowflake database. This issue has the highest priority, since grid planners do not want to wait for results.

Secondly, The current dashboard needs to be deployed within the Enexis landscape.

Additionally, some aesthetic improvements are welcome, since the presentation is quite basic at this moment.

DEPLOYMENT

7.1 Current Status

7.1.1 Code Standards

The python project is structured similarly as [Cookiecutter](#).

The code has been validated with linting by [Black](#), [isort](#), and [Flake8](#).

Unit tests have not been written yet.

7.1.2 Environment

The project currently uses data from the Production environment, since there is no DALI (meta)data available on the Test or Acceptance environments.

Therefore, the further implementation of the CICD pipeline is postponed for now. Also since within Enexis Gitlab is used, but for this study project a parallel Github repository is used.

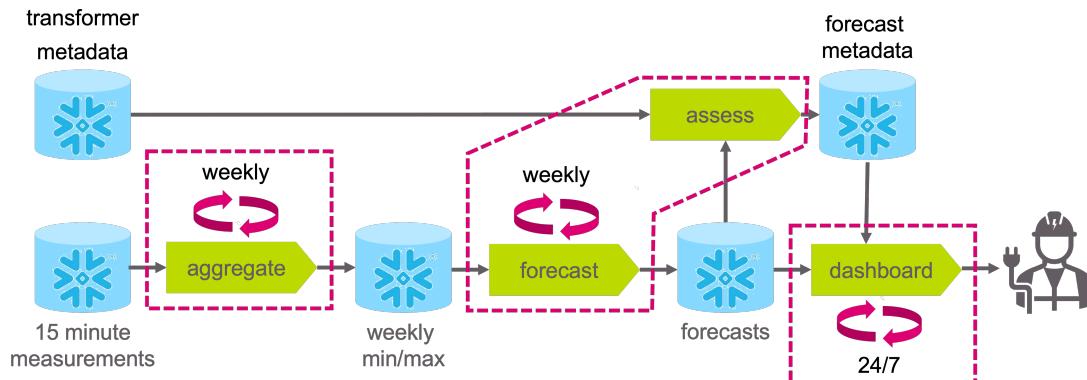
7.1.3 Security

Within the Enexis domain one logs on to an environment (in this case Production).

Credentials for accessing the Snowflake database are acquired via Vault to ensure to keep secrets secret.

7.1.4 Deployability

Docker files are written for following steps:



The steps are triggered via an entrypoint defined in the `docker-compose.yml`:

- `docker-compose run update_extremes` Aggregates the 15-minute data into weekly minimum and maximum values per DALI-box. Only updates new data since the last run.
- `docker-compose run make_forecasts` Creates a tuned model for a DALI box with (high) enough measurements and makes a forecast with it. Currently, the results are stored all with the `is_valid` column on `True`. An extra run should be done to set this column to `False` for old forecast. Until then, one should empty the table on beforehand.
- `docker-compose run -p 8000:8000 dashboard` Will run the dashboard and expose it to port 8000.

7.2 Deployment Plan

The following technical actions are required for full deployment:

- (Meta)data for a few DALI boxes is required on the Test and Acceptance environments.
- The project will need an app number for to trace the costs.
- The Gitlab CICD pipeline has to be configured.
- The project needs to land in the Artifactory repository manager of Enexis.
- Airflow DAGs have to be configured to trigger the jobs.

The end-users will eventually have to accept the product to go into Production. For that several end-user/stakeholder meetings are required to get valuable feedback and fine tune the product. During these meetings the project will be assessed on the success criteria.

Usage monitoring of the tool has not been implemented, is important after deployment for value evaluation and life cycle management.

7.3 Monitoring and Maintenance Plan

Monitoring of the triggered jobs will be done by connecting the Airflow job to a MS Team webhook. This automatically informs the Data Science Team if a scheduled job has failed.

The issue of model drift is obsolete, since every forecast is based on a newly trained model.

Monitoring of usage has to be implemented as mentioned above.

7.4 Final Report

One can create a pdf version of this documentation by the following steps:

- Clone the repository.
- Run `make latexpdf`

7.5 Project review

This section covers my personal growth reflection on the Lead track of JADS.

Firstly, I will cover the objectives and criteria set by JADS.

Secondly, I will review my personal goals.

7.5.1 JADS - Professional Education - Lead Track

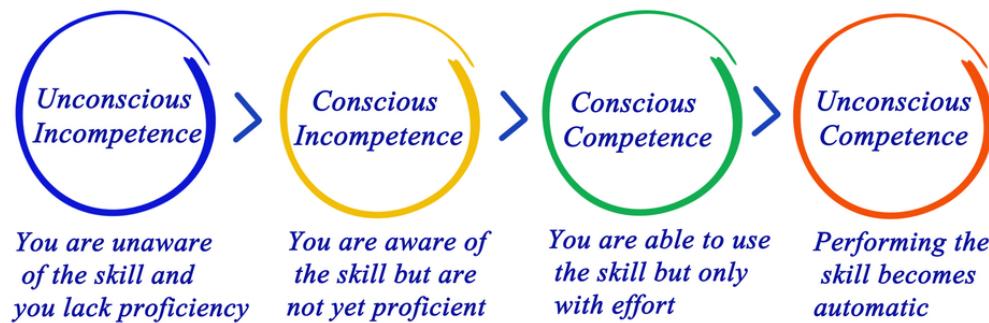
Learning Objectives

The general learning objectives for the Professional Education Lead track of JADS are:

- **Acquiring the knowledge and skills covered in the courses (technical and non-technical).** The most valuable courses were the non-technical ones. Partially, because there was maybe more to gain for me as a more technical character. But mainly, since they inspired me and enabled me to accept to embrace that skill set and also provided me with concrete tips to put it into practice.

Some technical lectures (e.g. “from POCs to Production” and “A Primer in Data Engineering”) confirmed that the way of working and view on data science within my professional team is on the right track. Other technical courses (e.g. “Bayesian Nets”, “Recent Advances in A.I.”) inspired me and sparked an interest in new subjects. Those subjects are now on my personal bucket list for a technical deep dive in the near future.

Concluding: I am convinced the courses in the Lead Track moved a good number of my skills towards the conscious incompetence stage.



The Conscious Competence Learning Model (source: pamelaslim.com).

- **Leading and implementing an impactful data science project by use of the CRISP-DM process.** The described project of this documentation shows the result of that.

- **Forming individuals to enable them to make impact with data science.** The coaching from the educators of JADS helped me reassess my future career.
 - The discussions with Jeroen about academic and engineering skills helped me appreciating my skills, my work and myself more.
 - The coaching of Kyril helped me to form a clearer view on my career goal and how to get there.
 - The highly effective presentation masterclass of Raoul helped me to reach goals more effective.

Evaluation Criteria

The evaluation criteria for the Professional Education Lead track of JADS and how I translated them are:

- Business Value.
 - Use case selection.
 - Stakeholder analysis.
 - Stakeholder and expectation management.
 - Value flow down diagram.
 - Regular business review moments.
- Programming.
 - Coding standards with Cookiecutter and linting (black, isort, flake8).
 - Using Sphinx for auto API documentation in GitHub.
 - Interactive visualisations in Altair.
 - Dashboard in Panel.
- Data engineering.
 - Snowflake database preprocessing usage (asynchronous queries, ETL).
 - Vault credential management.
 - Docker (compose with entrypoint) for every process step.
- Data analytics & machine learning.
 - Coding data management and model from scratch (since not using sklearn).
 - Probabilistic modelling in PyMC3.
 - Using Fourier and Taylor series in generalized additive model for time series forecast.
- Professional standard of reporting.
 - Using Sphinx for this documentation.
 - Status slide deck after sprint review with end users.
 - Project status / pitch and management summary for JADS peers and training.
- Academic / research skills.
 - Using the engineering approach (instead of the scientific method or the axiomatic system) to iteratively create and validate model and outcome.

Personal Goals

Additionally, there were also my personal learning objectives:

- Going through all the steps of CRISP-DM by myself (from Business Understanding to Deployment).
- Better understanding of probabilistic modeling techniques (e.g. bayesian models, probabilistic tools).
- Able to value my data skills and knowledge and creating traction / more confidence on flourishing as a data scientist.

One by one my personal learning objectives are fulfilled.

However, my conscious incompetence skill set has grown dramatically. So my new learning objectives are already there. (I will just have to order that backlog.)

To the persistent reader

This project is named Spark within Enexis.

But it has not just solely been named after the developer: Spa

API REFERENCE

This page contains auto-generated API reference documentation¹.

8.1 src

8.1.1 Subpackages

`src.forecast`

Submodules

`src.forecast.assess`

Module Contents

`src.forecast.assess.logger`

`src.forecast.assess.asses_forecasts (df_total, df_meta)`

Assesses if and when a transformer will overload based on the forecasts and metadata.

Parameters

`df_total: pd.DataFrame` DataFrame with forecasts.

`df_meta: pd.DataFrame` DataFrame with metadata of the transformer.

Returns

`pd.DataFrame` Result of the assessment.

¹ Created with sphinx-autoapi

`src.forecast.forecast`

Module Contents

`src.forecast.forecast.logger`

`src.forecast.forecast.config`

`src.forecast.forecast.determine_estimates(df_observed)`

Determine model parameters and forecast for a extreme (max or min).

Parameters

`df_observed` [pd.DataFrame] observed values of DALI data

Returns

`pd.DataFrame` estimates for extreme

`src.forecast.forecast.determine_estimates_minmax(df_observed)`

Determine model parameters and forecast for a extreme (max or min).

Parameters

`df_observed` [pd.DataFrame] observed values of DALI data

Returns

`pd.DataFrame` estimates for min and max

`src.forecast.forecast.forecast(boxid=None)`

Forecast for one or more DALI boxes.

Parameters

`boxid` [list] string or list with boxids to forecast for.

Returns

`pd.DataFrame` DataFrame with observations and forecast on long format

`src.forecast.make_forecasts`

Module Contents

`src.forecast.make_forecasts.logger`

`src.forecast.make_forecasts.make_forecasts(clear=False)`

Make forecasts for all boxes.

Parameters

`clear: bool` Clear the Snowflake table

Returns

`None`

`src.model`

Submodules

`src.model.format`

Module Contents

`src.model.format.logger`

`src.model.format.make_quantile_bands (df_base, samples, quantiles=5, 15, 50, 85, 95)`

Translate samples to bands with edges defined by the given quantiles and merge them to a base DataFrame.

Parameters

df_base [pd.DataFrame] base DataFrame with timestamps/length in accordance of the samples shape

samples [np.array] the posterior predictive samples

quantiles: list an iterable with the edges of the bands (0,1), ordered increasingly.

Returns

pd.DataFrame long format version with the quantile bands of the samples

`src.model.format.format_model_estimates (df_base, pp, quantiles=5, 15, 50, 85, 95)`

Format the samples into quantile bands for every model variable.

Parameters

df_base [pd.DataFrame] base DataFrame with timestamps/length in accordance of the samples shape

pp [dict] per model variable (key) the posterior predictive samples

quantiles: list an iterable with the edges of the bands (0,1), ordered increasingly.

Returns

pd.DataFrame long format version with the quantile bands of the samples for every model variable

`src.model.model`

Module Contents

`src.model.model.logger`

`src.model.model.det_dot (a, b)`

Dot product for Theano.

The theano dot product and NUTS sampler don't work with large matrices. Copyright (c) 2021. Ritchie Vink source: <https://www.ritchievink.com/blog/2018/10/09/build-facebook-prophet-in-pymc3-bayesian-time-series-analysis-with-generalized-additive-models/>

Parameters

a [np.array]

b [tt.vector]

Returns

np.array dot product of the two.

`src.model.model.fourier_series(t, p=52.1775, n=5)`

Calculate fourier representation of t for a period and order.

Copyright (c) 2021. Ritchie Vink Based on source: <https://www.ritchievink.com/blog/2018/10/09/build-facebook-prophet-in-pymc3-bayesian-time-series-analysis-with-generalized-additive-models/>

Parameters

t [range] range to be used as input variable

p [float] period to use for the fourier orders

n [int] order of fourier series

Returns

np.array matrix fourier representation of t

`src.model.model.seasonality_model(t, p=52.1775, n=5, seasonality_prior_scale=1)`

Create seasonality model with fourier series.

Copyright (c) 2021. Ritchie Vink Based on source: <https://www.ritchievink.com/blog/2018/10/09/build-facebook-prophet-in-pymc3-bayesian-time-series-analysis-with-generalized-additive-models/>

Parameters

t [range] range to be used as input variable

p [float] period to use for the fourier orders

n [int] order of fourier series

seasonality_prior_scale: float

Returns

pm.Var PYMC3 variable

`src.model.model.polynomial(t, n=4)`

Calculate polynomial representation of t for an order.

Parameters

t [range] range to be used as input variable

n [int] order of polynomial

Returns

np.array matrix polynomial representation of t

`src.model.model.drift_model(t, n=4)`

Polynomial drift/trend function for additive model.

Parameters

t [range] range to be used as input variable

n [int] order of polynomial.

Returns

pm.Var PYMC3 variable

```
src.model.model.create_model(t, y, p_fourier, n_fourier=5, n_polynomial=2)
Create a PYMC3 GAM model with a trend/drift and a seasonal/yearly component.
```

Parameters

t [timestamps] input series of scaled timestamps

y [float] observed values

p_fourier: float scaled period of the timestamps to take for the fourier component (a year)

n_fourier [int] order of the fourier component

n_polynomial: int order of the polynomial component

Returns

PYMC3 model context

```
src.plot
```

Submodules

```
src.plot.altair
```

Module Contents

```
src.plot.altair.plot_estimate(df, legend=True)
```

Plot the load forecast transformer.

Parameters

df [pd.DataFrame] DataFrame with the columns date, lower, upper, forecast(Q10-Q190, median), extreme

Returns

Altair chart

```
src.plot.altair.plot_observed(df)
```

Plot the historic/observed load of a transformer.

Parameters

df [pd.DataFrame] DataFrame with like:

Returns

Altair chart

```
src.plot.altair.plot_limits(df)
```

Plot the capacity limits of a transformer as a ruler of as a line if metadata changes over time.

Parameters

df [pd.DataFrame] DataFrame with the limits to be plotted

Returns

Altair chart

```
src.plot.altair.lightness_scale(factor, limits=['#bedef4', '#1f77b4'])
```

Create right color intensity for float values.

Parameters

factor: float Factor of full color intensity.
limits: list Hex colour codes for start and end of scale.

Returns

str Hex code for the factor value.

`src.plot.altair.plot_base(df_data=None, df_meta=None)`
Plot one or more of the following: history, forecast, transformer limits.

Parameters

df_data [pd.DataFrame] DataFrame with the timeseries
df_meta [pd.DataFrame] DataFrame with the metadata

Returns

Altair layer chart

`src.plot.altair.plot_decompose(df)`
Plot decomposition of PYMC3 GAM model components by model variable.

Parameters

df [pd.DataFrame] Data of estimates in long format.

Returns

alt.Chart Altair plot.

`src.plot.altair.plot_all(df_data, df_meta=None)`
Plot estimates, observed and limits.

Parameters

df_data [pd.DataFrame] DataFrame with the timeseries
df_meta [pd.DataFrame] DataFrame with the metadata

Returns

Altair layer chart

`src.plot.dashboard`

Module Contents

`src.plot.dashboard.build_dashboard()`
Initialize dashboard.

Returns

None

`src.plot.dashboard.dashboard`
`src.plot.dashboard.dashboard`

```
src.plot.format
```

Module Contents

```
src.plot.format.format_limits(df_meta, df_data=None)
```

Format the transformer limits for plotting.

Parameters

df_meta: pd.DataFrame DataFrame with the transformer limits.

df_data: pd.DataFrame Optionally, the historic preprocess to determine the time range for the limits.

Returns

A DataFrame with the limits to plot.

```
src.plot.format.format_history(df_data)
```

Format the historic for plotting.

Parameters

df_data: pd.DataFrame The historic preprocess to be plotted.

Returns

A DataFrame with the historic preprocess to plot.

```
src.plot.format.format_forecast(df_forecast)
```

Format the forecasted load.

Parameters

df: pd.DataFrame

Returns

pass through input

```
src.plot.format.dummy_forecast(df)
```

Create a dummy forecast with median and Q10-Q90 based on historic preprocess to test plotting function.

Parameters

df [pd.DataFrame] Historic preprocess to use as a basis

Returns

pd.DataFrame Forecast preprocess that can be used to plot with plot_forecast()

```
src.preprocess
```

Submodules

```
src.preprocess.preprocess
```

Module Contents

```
src.preprocess.preprocess.logger
```

```
src.preprocess.preprocess.too_short(df_data, threshold=52)
Check if number of entries is long enough.
```

Parameters

df_data: pd.DataFrame Data to check.

threshold: int Minimum lenght.

Returns

bool True if preprocess is long enough.

```
src.preprocess.preprocess.too_small(df_data, capacity, threshold=0.25)
Check if values of preprocess is greater than threshold x capacity of transformer.
```

Parameters

df_data: pd.DataFrame Data to check.

capacity: int Capacity of the transformer.

threshold: int Fraction of the capacity that should be reached.

Returns

bool True if preprocess exists that is greater than threshold x capacity.

```
src.preprocess.preprocess.remove_leading_idling(df_data, capacity, threshold=0.01)
Remove preprocess that was generated when DALI box was active but electrical connection was not.
```

Parameters

df_data [ps.DataFrame] Data to be cleaned.

capacity: int Capacity of the transformer.

threshold: int Fraction of the capacity that should be reached.

Returns

pd.DataFrame DataFrame without the idling preprocess points.

```
src.preprocess.preprocess.load_data(boxid)
Load preprocess and metadata for boxid and do preprocessing.
```

Parameters

boxid: str ID of DALI box.

Returns

None | (**df_data**, **df_meta**) If checks are OK, return DataFrames with historic and meta pre-process.

```
src.preprocess.preprocess.format_data(df)
Format the loaded weekly data in the correct form.
```

Parameters

df: pd.DataFrame Loaded weekly extremes.

Returns

pd.DataFrame Formatted data.

```
src.preprocess.preprocess.split_last(df_data, period=dt.timedelta(weeks=26))
Split the historic dataset into a train and test set a certain period from the end.
```

Parameters

df_data: pd.DataFrame Dataset to split.
period: datetime Period from the end to split from

Returns

—— df_train, df_test

```
src.preprocess.preprocess.extrapolate_timestamps(df, horizon=dt.timedelta(weeks=26))  
Extrapolate the data with timestamps for the future.
```

Parameters

df [pd.DataFrame] data to extrapolate.
horizon [dt.timedelta] the amount of time to extrapolate further.

Returns

pd.DataFrame Only the extrapolated/added part.

```
src.preprocess.update_extremes
```

```
src.utils
```

Submodules

```
src.utils.logger
```

Module Contents

```
src.utils.logger.init_logging()  
Initialize logger.
```

Returns

logger object

```
src.utils.parser
```

Module Contents

```
src.utils.parser.parse_config(abs_path) → EnvYAML  
Parse the configuration in EnvYAML format.
```

```
src.utils.parser.parse_args() → dict  
Parse the command-line arguments.
```

`src.utils.preprocess`

Module Contents

`src.utils.preprocess.downcast(s, try_numeric=True, category=False)`

Downcast a series to the lowest possible memory type.

Parameters

`s` [pd.Series] Series to downcast.

`try_numeric: bool` If True it will try to read strings as numeric values.

`category: bool` If True (string) objects will be cast as a category.

Returns

`Downcasted series.`

`src.utils.preprocess.map_labels(series, kind='categorical', labels=None, backwards=False, **arg)`

Map a Series values by the labels given.

Parameters

`series: pd.Series` Series to map on.

`kind: str` Indicator for kind of preprocess in series. With kind of {"categorical", "ordinal"} the mapping is applied, otherwise not.

`labels: dict` Defines with the mapping {key_0: value_0, etc.}.

`arg:` Additional arguments.

Returns

`pd.Series` Series with mapped values.

`class src.utils.preprocess.MinMaxScaler(upper=1, lower=-1)`

MinMax Scaler like in sklearn, prevents total library import/dependency.

Initialize scaler with upper and lower boundary.

Parameters

`upper` [float] upper boundary to scale to

`lower` [float] lower boundary to scale to

`fit(self, X, y=None)`

Get fit parameters.

Parameters

`X` [np.array] preprocess to fit on

`y` [None] solely for consistency

Returns

`self` instance with self.min, self.max defined.

`transform(self, X)`

Scales preprocess according to fitted parameters.

Parameters

X [np.array] preprocess to scale

Returns

np.array scaled preprocess

fit_transform(self, X, y=None)

Execute consecutively self.fit and self.transform.

Parameters

X [np.array] preprocess to scale

y [None] solely for consistency

Returns

np.array scaled preprocess

inverse_transform(self, X, y=None)

Scale back to original domain.

Parameters

X [np.array] preprocess to scale

y [None] solely for consistency

Returns

np.array scaled preprocess

src.utils.snowflake

Module Contents

src.utils.snowflake.logger

src.utils.snowflake.config

src.utils.snowflake.channel_like = register://electricity/0/activepower/?avg=15

src.utils.snowflake.column_details =

(BOXID VARCHAR(50), L VARCHAR(5), YEAR NUMBER(4), WEEK NUMBER(2), PROCESSED_ON TIMESTAMP_TZ, MAX DOUBLE, MIN DOUBLE)

src.utils.snowflake.format_connection(name)

Get (connection) details in the right format for different tables.

Parameters

name [name of the different sources]

Returns

Connection

src.utils.snowflake.read_meta(boxid=None)

Read meta preprocess of DALI box.

Parameters

boxid [list] ID of DALI box to read. None results in all available preprocess of DALI boxes that have nominal power specified.

Returns

pd.DataFrame with metadata.

```
src.utils.snowflake.make_week_extremes_query(boxid=None,           last_processed=dt.datetime(2001, 1, 1))
```

Build the query to request week extremes.

Parameters

boxid [list] If not None: select specific boxes.

last_processed: datetime Date to determine week extremes from. Cuts are always made on the last monday.

Returns

str Query string.

```
src.utils.snowflake.create_table_query(query)
```

Make query to create or replace table and insert preprocess of select query.

Parameters

query [str] Select query to be used.

Returns

str New query

```
src.utils.snowflake.insert_table_query(query)
```

Make query to insert preprocess of select query in an existing table.

Parameters

query [str] Select query to be used.

Returns

str New query

```
src.utils.snowflake.create_week_extremes()
```

Execute query to create whole new table of week extremes asynchronously.

```
src.utils.snowflake.get_last_processed_time()
```

Retrieve last time table update has been done.

Returns

datetime Last processing time.

```
src.utils.snowflake.update_week_extremes()
```

Update week extremes from last processing time up to last monday.

```
src.utils.snowflake.read_week_extremes(boxid=None, L=None)
```

Read week extremes for a DALI box and phase.

Parameters

boxid: list If not None: A list with DALI box IDs to read.

L: str The phases to retrieve (sumli, L1, L2, L3)

Returns

pd.DataFrame Week extremes.

```
src.utils.snowflake.clear_forecasts()
```

Clear the table with forecasts.

Returns**None**`src.utils.snowflake.clear_forecast_meta()`

Clear the table with forecast metadata.

Returns**None**`src.utils.snowflake.write_forecasts(df)`

Write results into the Snowflake database.

Parameters**df: pd.DataFrame** Results to write.**Returns****None**`src.utils.snowflake.write_forecast_meta(df)`

Write assessment of forecast on capacity to forecast metadata table in Snowflake.

Parameters**df: pd.DataFrame** Assessment results.**Returns****None**`src.utils.snowflake.read_forecasts(boxid=None)`

Read forecasts for Snowflake database.

Parameters**boxid: str** Boxid of a DALI box. If not provided or None, all forecast are loaded.**Returns****pd.DataFrame** DataFrame with forecasts.`src.utils.snowflake.read_forecast_meta()`

Read assessment of forecasts.

Returns**pd.DataFrame** Results of he assessment of all DALI box forecasts.`src.utils.snowflake.get_forecasted_boxids()`

Get the boxids for boxes that are already forecasted.

(Should be adapted when multiple forecasts (done on different dates) are present in database).

Returns**pd.DataFrame** Unique boxids.

`src.utils.vault`

Module Contents

`src.utils.vault.get_secrets(connection)`

Get secrets from Vault.

Parameters

`connection: str` The application/connection to get secrets for.

Returns

`requested secrets`

`src.utils.vault.get_vault_secret(url, token, path, mount_point, **kwargs)`

Get vault secrets to certain path in a dictionary.

Parameters

`url [str]` Base URL of Vault server

`token [str]` Token to get access to Vault

`path [str]` Relative path to location of secret

`mount_point [str]` Mount point of address

Returns

`dict` Dictionary with keys and value of a secret

8.1.2 Package Contents

`src.init_logging()`

Initialize logger.

Returns

`logger object`

PYTHON MODULE INDEX

S

src, 43
src.forecast, 43
src.forecast.assess, 43
src.forecast.forecast, 44
src.forecast.make_forecasts, 44
src.model, 45
src.model.format, 45
src.model.model, 45
src.plot, 47
src.plot.altair, 47
src.plot.dashboard, 48
src.plot.format, 49
src.preprocess, 49
src.preprocess.preprocess, 49
src.preprocess.update_extremes, 51
src.utils, 51
src.utils.logger, 51
src.utils.parser, 51
src.utils.preprocess, 52
src.utils.snowflake, 53
src.utils.vault, 56

INDEX

A

asses_forecasts() (in module `src.forecast.assess`), 43

B

build_dashboard() (in module `src.plot.dashboard`), 48

C

channel_like (in module `src.utils.snowflake`), 53
clear_forecast_meta() (in module `src.utils.snowflake`), 55
clear_forecasts() (in module `src.utils.snowflake`), 54
column_details (in module `src.utils.snowflake`), 53
config (in module `src.forecast.forecast`), 44
config (in module `src.utils.snowflake`), 53
create_model() (in module `src.model.model`), 46
create_table_query() (in module `src.utils.snowflake`), 54
create_week_extremes() (in module `src.utils.snowflake`), 54

D

dashboard (in module `src.plot.dashboard`), 48
det_dot() (in module `src.model.model`), 45
determine_estimates() (in module `src.forecast.forecast`), 44
determine_estimates_minmax() (in module `src.forecast.forecast`), 44
downcast() (in module `src.utils.preprocessing`), 52
drift_model() (in module `src.model.model`), 46
dummy_forecast() (in module `src.plot.format`), 49

E

extrapolate_timestamps() (in module `src.preprocessing.preprocessing`), 51

F

fit() (`src.utils.preprocessing.MinMaxScaler` method), 52
fit_transform() (`src.utils.preprocessing.MinMaxScaler` method), 53

forecast() (in module `src.forecast.forecast`), 44
format_connection() (in module `src.utils.snowflake`), 53
format_data() (in module `src.preprocessing.preprocessing`), 50
format_forecast() (in module `src.plot.format`), 49
format_history() (in module `src.plot.format`), 49
format_limits() (in module `src.plot.format`), 49
format_model_estimates() (in module `src.model.format`), 45
fourier_series() (in module `src.model.model`), 46

G

get_forecasted_boxids() (in module `src.utils.snowflake`), 55
get_last_processed_time() (in module `src.utils.snowflake`), 54
get_secrets() (in module `src.utils.vault`), 56
get_vault_secret() (in module `src.utils.vault`), 56

I

init_logging() (in module `src`), 56
init_logging() (in module `src.utils.logger`), 51
insert_table_query() (in module `src.utils.snowflake`), 54
inverse_transform() (`src.utils.preprocessing.MinMaxScaler` method), 53

L

lightness_scale() (in module `src.plot.altair`), 47
load_data() (in module `src.preprocessing.preprocessing`), 50
logger (in module `src.forecast.assess`), 43
logger (in module `src.forecast.forecast`), 44
logger (in module `src.forecast.make_forecasts`), 44
logger (in module `src.model.format`), 45
logger (in module `src.model.model`), 45
logger (in module `src.preprocessing.preprocessing`), 49
logger (in module `src.utils.snowflake`), 53

M

make_forecasts() (in module `src.forecast.make_forecasts`), 44

```
make_quantile_bands () (in module src.model.format), 45
make_week_extremes_query () (in module src.utils.snowflake), 54
map_labels () (in module src.utils.preprocess), 52
MinMaxScaler (class in src.utils.preprocess), 52
module
    src, 43
    src.forecast, 43
    src.forecast.assess, 43
    src.forecast.forecast, 44
    src.forecast.make_forecasts, 44
    src.model, 45
    src.model.format, 45
    src.model.model, 45
    src.plot, 47
    src.plot.altair, 47
    src.plot.dashboard, 48
    src.plot.format, 49
    src.preprocess, 49
    src.preprocess.preprocess, 49
    src.preprocess.update_extremes, 51
src.utils, 51
src.utils.logger, 51
src.utils.parser, 51
src.utils.preprocess, 52
src.utils.snowflake, 53
src.utils.vault, 56
```

P

```
parse_args () (in module src.utils.parser), 51
parse_config () (in module src.utils.parser), 51
plot_all () (in module src.plot.altair), 48
plot_base () (in module src.plot.altair), 48
plot_decompose () (in module src.plot.altair), 48
plot_estimate () (in module src.plot.altair), 47
plot_limits () (in module src.plot.altair), 47
plot_observed () (in module src.plot.altair), 47
polynomial () (in module src.model.model), 46
```

R

```
read_forecast_meta () (in module src.utils.snowflake), 55
read_forecasts () (in module src.utils.snowflake), 55
read_meta () (in module src.utils.snowflake), 53
read_week_extremes () (in module src.utils.snowflake), 54
remove_leading_idling () (in module src.preprocess.preprocess), 50
```

S

```
seasonality_model () (in module src.model.model), 46
```

```
split_last () (in module src.preprocess.preprocess), 50
src
    module, 43
    src.forecast
        module, 43
    src.forecast.assess
        module, 43
    src.forecast.forecast
        module, 44
    src.forecast.make_forecasts
        module, 44
src.model
    module, 45
src.model.format
    module, 45
src.model.model
    module, 45
src.plot
    module, 47
src.plot.altair
    module, 47
src.plot.dashboard
    module, 48
src.plot.format
    module, 49
src.preprocess
    module, 49
src.preprocess.preprocess
    module, 49
src.preprocess.update_extremes
    module, 51
src.utils
    module, 51
src.utils.logger
    module, 51
src.utils.parser
    module, 51
src.utils.preprocess
    module, 52
src.utils.snowflake
    module, 53
src.utils.vault
    module, 56
```

T

```
too_short () (in module src.preprocess.preprocess), 49
too_small () (in module src.preprocess.preprocess), 50
transform () (src.utils.preprocess.MinMaxScaler method), 52
```

U

```
update_week_extremes () (in module src.utils.snowflake), 54
```

W

```
write_forecast_meta()      (in      module  
    src.utils.snowflake), 55  
write_forecasts() (in module src.utils.snowflake),  
    55
```